

#### Question:

learning\_rate: learning rate to be used for all updates, defaults to 0.001

momentum: momentum/gamma values for momentum based gradient descent, defaults to 0.5

num\_hidden: number of hidden layers

sizes: comma separated values, number of units in each hidden layer

activation: non-linearity to be used after each layer, defaults to relu

loss: loss function to optimize, defaults to "ce" i.e cross entropy

opt: optimizer to be used for learning, defaults to "adam". Could be any one of "gd", "momentum", "adam", "nag"

batch\_size: size of minibatch, default is 20

num\_epochs: number of epochs to run for, defaults to 10

anneal: a boolean argument. If set to "true", validation loss of the current epoch is compared with the previous epoch and if found to be greater than the latter, the epoch is restarted and any weight updates during the epoch are discarded.

path\_save\_dir: path to the folder where the final model is stored

expt\_dir: path to the folder where the logs are stored

path\_to\_train: path to the training data .csv file

path\_to\_val: path to the validation dataset .csv file

path\_to\_test: path to the test dataset .csv file

to\_pretrain: boolean variable, if set to "true", weights from the epoch determined by state from a previous training session are loaded

epoch\_to\_restore: epoch from which weights need to be loaded

to\_test: boolean variable, if set to "true", weights determined by "state" are loaded and predictions produced for the test data

#### Notes:

Used Xavier initialization for weights so that the variance of the outputs of each layer is equal to the variance of the inputs it helps to prevent the gradients from vanishing or exploding during training.

Implemented the following optimizers:

1. Gradient Descent
2. Momentum
3. Adam
4. Nesterov Accelerated Gradient

Took the data from scikit-learn datasets